

31251 – Data Structures and Algorithms

Week 5 - Graphs Part I

Luke Mathieson

Does anyone read these titles?

- Graphs
- Traversing Graphs
- Greedy Algorithms
 - Examples: Prim's Algorithm, Kruskal's Algorithm

Graphs

Graphs as Mathematical Objects

- Graphs are an incredibly useful modelling tool.
- Simple graphs consist of:
 - A set of elements called *vertices*.
 - A set of unordered pairs of distinct vertices called *edges*.
 - There is only one edge between a pair of vertices.
- Other types of graph come from altering these conditions:
 - Ordered pairs gives *directed* graphs.
 - More than one edge per pair gives *multi-graphs*.
 - More than two vertices per edge gives *hypergraphs*.
 - Edges (and vertices) can be weighted, labelled, &c.

Graphs as Mathematical Objects

- The vertices model interesting things:
 - Computers
 - Processes
 - Proteins
 - Production facilities
 - Websites
- The edges model relationships between interesting things:
 - Network links
 - Shared resources
 - Protein interactions
 - Transport links
 - Hyper links
- Used somehow in virtually every part of computer science.

Some Notation and Definitions

- If two vertices have an edge between them, they are *adjacent*.
- If a vertex is one of the pair that forms an edge, it is *incident* to that edge.
- The number of edges incident to a vertex is the *degree* of that vertex.
- Graphs will be denoted with uppercase letters like G , H , &c.
- Vertices will be denoted by lowercase letters like u , v , &c., or natural numbers $1, 2, 3, 4, \dots, n$.
- Edges will be denoted by lowercase letters like e , d , &c. or by their incident vertices:
 - In undirected graphs uv .
 - In directed graphs (u, v) – u is the *tail*, v is the *head*.
 - Directed edges are also called *arcs*.

- As an abstract data structure, a graph needs to support a lot of basic operations:
 - `addVertex(Vertex v)`
 - `removeVertex(Vertex v)`
 - `addEdge(Vertex u, Vertex v)`
 - `removeEdge(Vertex u, Vertex v)`
 - `adjacent(Vertex u, Vertex v)`
 - `degree(Vertex u)`
 - return the vertices in the graph
 - return the edges incident to a vertex
 - return the vertices adjacent to a vertex
- As well as the normal things needed for a usable implementation in a given language (constructors, &c.)
- Some operations may depend on the exact implementation.

Graphs as Data Structures – Adjacency Matrix

- Simplest form:
 - Edges are stored as a two-dimensional matrix (e.g. `vector<vector<bool> > edges` or `bool edges[] []`).
 - `edge[i][j] == true` means vertex i is adjacent to j .
- Some enhancements:
 - Can use a numeric (`int`, `double`, ...) matrix to give weighted edges.
 - Can use a matrix of Edges if edges are more complex – a null-type value means no edge.
 - Easily supports directed graphs and undirected graphs.
 - If vertices have associated data, we can store them separately (another array would make matching indices easy).
- Quick access – $O(1)$, not so great space and set-up – $O(n^2)$ (with n vertices).

Graphs as Data Structures – Adjacency List

- Each vertex has associated with it a list of its adjacent vertices.
- Could be a size n array of linked lists, or similar.
- Slower to determine adjacency – $O(n)$, but faster to return all adjacent vertices – $O(1)$.
- Most compact space representation $O(m + n)$ where m is the number of edges in the graph – we have to store something for each vertex and edge anyway, so this is the best we can do.
- Easy to modify for more complex edge and vertex data structures.
- Works best for sparse graphs (few edges per vertex).

Graphs as Data Structures – Object Oriented

The extreme version:

- We have classes for Vertex, Edge and Graph.
- Vertex contains a list of its Edges.
- Each Edge knows its endpoints.
- the Graph knows about everything.
- *Lots* of references to keep track of.
- Tends to be the slow way to do things, but has a nice match to the conceptual version.

Traversing Graphs

Moving About in Graphs

- Many algorithms rely on being able to explore the graph.
- We need to be able to keep track of which vertices we've been to.
- We also need a way of picking which neighbour to move to next:
 - We can use an inherent order on the vertices (by label, number, &c.), or pick an arbitrary order.

- Pick a starting vertex.
- Recursively pick an unvisited neighbour and visit it.
- Can be implemented recursively, or iteratively using a stack.

Depth First Traversal – Recursive

```
dft(Vertex v){  
    mark v as visited;  
    visit(v);  
    for each neighbour u of v{  
        if (u is unmarked)  
            dft(u);  
    }  
}
```

Depth First Traversal – Iterative

```
dft(){
    pick starting vertex v;
    Stack unprocessed = new Stack();
    unprocessed.push(v);

    while (!unprocessed.isEmpty()){
        Vertex u = unprocessed.pop();
        if (u is unmarked){
            visit(u);
            mark u;
            for each neighbour w of u{
                unprocessed.push(w);
            }
        }
    }
}
```

- Pick a starting vertex and put it in a queue.
- Iteratively take a vertex from the queue, visit it and place all its neighbours in the queue.
- It's inherently iterative, it's not impossible to implement recursively, just silly.

Breadth First Traversal – Iterative

```
bft(){
    pick starting vertex v;
    Queue unprocessed = new Queue();
    unprocessed.offer(v);

    while (!unprocessed.isEmpty()){
        Vertex u = unprocessed.poll();
        if (u is unmarked){
            visit(u);
            mark u;
            for each neighbour w of u{
                unprocessed.offer(w);
            }
        }
    }
}
```

Breadth First vs. Depth First

- Some graphs produce the same traversal order for both.
- Which one to use will depend upon the application.
- Notice the iterative versions of both are actually identical – just swap the stack and the queue.
- Both $O(n + m)$. (Why?)

Greedy Algorithms

- Greedy Algorithms are based on picking what looks good now.
- One of the simplest algorithmic paradigms.
- Usually easy to implement.
- Only really works for certain types of problems.
- For some problems, a greedy approach may produce the worst solution!

When do Greedy Algorithms work?

- Greedy Algorithms usually work when the problem satisfies two properties:
 - **Optimal Substructure:** Optimal solutions contain optimal subsolutions to subproblems (you can split solutions up and they're still good).
 - **Greedy Choice Property:** Any decisions depend only on what you've already seen (you don't have to come back and fix things).

Spanning Trees of Graphs

Consider the following problem:

- A company has to connect cities with fibre optic cable such that each city has a (possibly multi-hop) link to every other city. The company knows the cost of linking each pair of cities, and wants to accomplish its task with a minimum cost.
- We can use a weighted graph to model this problem, but what are we looking for?
 - A set of edges that connects all the vertices.
 - No unneeded edges.
 - It's a thing called a tree! (See next slide)
- So we want a tree that includes all the vertices, and has the minimum total edge weight.

- A *subgraph* is a subset of the vertices and edges of a graph that form a graph.
- A subgraph is *spanning* if it includes all the vertices of the original graph.
- A spanning subgraph is a *spanning tree* if it contains no cycles.
- A spanning tree is a *minimum spanning tree* if it has minimum total (edge) weight over all possible spanning trees of that graph (is it unique?).

- In unweighted graphs (or a graph where all edge weights are the same), *any* spanning tree is a minimum spanning tree.
- We can compute one from a depth-first or breadth-first traversal.

```
df_spanning_tree(Vertex v, Tree t){  
    mark v as visited;  
    for each neighbour u of v{  
        if (u is unmarked)  
            add edge vu to t;  
            dft(u,t);  
    }  
}
```

- If we have different weights on the edges, a simple traversal is not enough.
- There are two main algorithms:
 - Prim's Algorithm
 - Kruskal's Algorithm
- If we have time, we'll look at Borůvka's (Sollin's) Algorithm.
- These are all greedy algorithms, with similar but slightly different approaches.

- Given a *connected* graph G
 - ① Pick a starting vertex v (however you want), add v to the partially complete tree T .
 - ② While $|T| < n$
 - ① Let E' be the set of edges uv where $u \in T$ and $v \in G \setminus T$.
 - ② Let uv be the edge of smallest weight in E' .
 - ③ Add uv to the edges of T and v to the vertices of T .
 - ③ Return T .

- In other words:
- Keep track of which edges and vertices are in the tree.
- Pick the next smallest edge that extends the tree, add it.
- Keep going until the whole graph is spanned.

Prim's Algorithm Complexity

- If we use an adjacency matrix, and search for the edges: $O(n^2)$.
- Putting the edges into a binary heap, with the graph stored as an adjacent list: $O((n + m) \log n) = O(m \log n)$.
- Using Fibonacci heap (don't worry about what this is) and adjacency lists: $O(m + n \log n)$.

- Prim's algorithm grows the spanning tree by greedily picking the best next edge.
- Kruskal's algorithm approaches the problem more globally - start with a lot of trees (a *forest*), and pick the best edge to connect two components.

- Given a *connected* graph G , start with n trees $\{T_i\}$, each with one vertex.
 - ① While there is more than one tree
 - ① Pick the smallest edge uv such that u is in one tree T_i , and v is in another T_j .
 - ② Merge T_i and T_j by adding uv .
 - ② Return the final tree T .

Kruskal's Algorithm Complexity

- By labelling vertices with which component they're in, we can get $O(n \cdot m)$ – not great.
- If we sorted the edges, and employ an efficient disjoint set data structure (haven't seen one in the course):
 $O(m \log m) = O(m \log n)$ – about the same as the normal Prim implementation.
- If the edges can be sorted efficiently by counting sort or radix sort or similar, we can get $O(m \cdot \alpha(n))$, where $\alpha(n)$ is the inverse of the single valued Ackermann function (look it up some time).

Extra Stuff on Spanning Trees

- Invented in 1926 by Otakar Borůvka – see computers aren't necessary for *computer science*.
- Reinvented three more times, lastly by Sollin in 1965.
- Works like a cross between Prim's and Kruskal's algorithms

- Given a graph G , initialise n tree $\{T_i\}$, each containing one vertex.
- ① While there is more than one tree
 - ① For each component tree
 - ① Pick the smallest outgoing edge (connecting this component to another)
 - ② Add this edge to the trees, merging them.
- ② Return the single remaining tree T

Borůvka's Algorithm – Complexity

- The outer loop only needs to execute $O(\log n)$ times – we halve the number of components at each step.
- Along with search for the edges at each iteration, we get $O(m \log n)$ without too much fiddling.
- A similar approach as used for Kruskal's can be used to get $O(m \cdot \alpha(n))$.
- A randomised version exists with $O(m)$ expected running time – remember this is linear in the size of the graph, about as fast as possible.

Correctness of Prim's Algorithm I

Lemma

Given a connected, weighted graph G , Prim's algorithm produces a minimum spanning tree of G .

Proof:

- As G is connected, it is (or at least should be) clear that Prim's algorithm produces a tree that spans the graph. Thus we need only argue that it is a *minimum* spanning tree.
- Let T_P be the tree produced by Prim's algorithm.
- Assume for contradiction that there exists a minimum spanning tree T_M of G and that the weight of T_M is less than the weight of T_P .
- Let e be the first edge added to T_P that is not in T_M , and let $S \subset V$ be the vertices in the partial tree at the point e is added.
- Note that e has one endpoint in S and the other in $V \setminus S$.

Correctness of Prim's Algorithm II

- As T_M is a spanning tree, there must be a path in the tree between the endpoints of e .
- On this path there must be some edge f in T_M with one endpoint in S and the other not.
- Then at the point of adding e , f must've been a candidate edge, that the algorithm didn't pick, hence the weight of f is at least the weight of e .
- If the weight of f is strictly greater than that of e , we can construct a new tree $T_{M'} = T_M - f + e$ with smaller weight than T_M , contradicting the assumption that T_M is minimum.
- If the weight of f is the same as that of e , we can construct $T_{M'}$, then repeat the argument with $T_{M'}$ in place of T_M – either we get a contradiction as before, or we progressively modify T_M to be T_P and therefore T_P must also be minimum.

Correctness of Kruskal's Algorithm I

Lemma

Given a connected, weighted graph G , Kruskal's algorithm produces a minimum spanning tree of G .

Proof:

- This time we use an inductive proof.
- What we will show is that if F is the set of edges chosen at any point in the algorithm, then there is some minimum spanning tree that contains F .

① Base Case:

- $F = \emptyset$. The proposition is trivially true in this case as \emptyset is a subset of any set.

② Inductive Assumption:

- Assume the algorithm (to this point) has produced edge set F' , and F' can be extended to some MST T .

③ Inductive Step:

Correctness of Kruskal's Algorithm II

- If the next chosen edge e is also in T , then our proposition holds for $F + e$.
 - If it is not, then $T + e$ contains a cycle, and there is some edge f that is in the cycle, but not in F .
 - The weight of f must be at least the weight of e – otherwise the algorithm would choose f at this point.
 - Then $T - f + e$ is a minimum spanning tree that contains $F + e$, and we're done.
- Then by induction the final set of edges can be 'extended' to an MST – as it spans the graph, this extension is 'do nothing', so the algorithm produces an MST.

Lemma

If e is the unique smallest weight edge in a connected, weighted graph G , then e is in every MST of G .

Proof: Assume for contradiction there is some MST T that does not contain e , then the graph $T + e$ contains a cycle, we can pick any edge from this cycle (other than e) and remove it to obtain a new spanning tree. As e has weight less than every other edge, this new tree must have smaller weight than T , therefore T was not an MST.

What happens if there's more than one minimum weight edge – are they all in every MST?

This can be inductively extended:

Lemma

If G is a connected, weighted graph where all edge weights are distinct, G has a unique minimum spanning tree.

And similar proofs give (assume G is connected):

Lemma

Let C be a cycle in G , and e be the largest weight edge in C . e is not in any MST of G .

Lemma

Let D be any cut in G , and e be the smallest weight edge in D . e is in every MST of G .

What about disconnected graphs?

- If the graph has several disconnected components, we can't get a spanning tree (trees have to be connected).
- We can get a spanning forest.
- The algorithms we have seen so far won't work (why not? can they be fixed?).

Reverse-Delete – Kruskal's Other Algorithm

- Appears in the same paper as Kruskal's algorithm.
- Sort of like a backwards Kruskal's.
- We remove edges, instead of adding them, and see what we're left with at the end.

Reverse-Delete – Kruskal's Other Algorithm

Given a weighted graph G :

- 1 Sort the edges by decreasing weight.
- 2 While edges remain to be processed
 - 1 Take the next biggest edge e .
 - 2 Check if deleting e will create more components than you already have.
 - 3 If not, delete it, otherwise keep it.
- 3 Return the remaining graph.

- We can sort the edges in $O(m \log m)$.
- We can check the connectivity in $O(\log n (\log \log n)^3)$.
- So in total we get $O(m \log n (\log \log n)^3)$ time.