

31251 – Data Structures and Algorithms

Week 9 - Graphs Part II

Luke Mathieson

Does anyone read these titles?

- Dynamic Programming
 - Example: Dijkstra's Algorithm
- Connectivity in Graphs
 - Connectivity and Connected Components
 - Strongly Connected Components
 - Articulation Points
- Dependencies in Directed Graphs

Dynamic Programming

- Dynamic programming looks a lot like Divide-and-Conquer, but subinstances *overlap*.
- Usually implemented with some sort of table that keeps track of subinstance solutions.
- Normally they work bottom-up (the reverse of D&C):
 - Solve the smallest subinstances, combine them to make bigger subinstances, keep going until you have the whole thing.
- Can be done top-down (recursively) as well.

What makes it work?

- **Optimal Substructure:** The same as greedy algorithms and D&C.
- **Overlapping Subproblems:** As mentioned before.
- That's it actually - it's a very general technique.

Shortest Paths

Single Source Shortest Paths

- In unweighted graphs, the distance between two vertices is easy to calculate:
 - Just do a breadth first search.
 - If you're looking for a particular vertex, stop when you find it.
 - The distance is the same as the number of iterations taken to get to the vertex.
- In weighted graphs, it's not so easy.
 - The lowest weight path is not always the path with fewest edges.

- The backbone algorithm for most shortest path algorithms.
- Employs a dynamic programming approach.
- It's not too far removed from Prim's algorithm
 - It builds a shortest path tree instead,
 - but a basic greedy approach doesn't work in deciding which edge to add.
- Doesn't handle negative edge weights (in the basic version).

Given a graph G with non-negative edge weights, with a given starting vertex u :

- ① Set the tentative distances for all vertices; 0 for u , infinity for all others.
- ② Mark all vertices as unvisited.
- ③ Set u as the current vertex.
- ④ While there are unvisited vertices:
 - ① For each unvisited neighbour of the current vertex:
 - ① Compare the tentative distance to the neighbour with the distance to current plus the edge weight of the edge joining them.
 - ② Keep whichever is smaller.
 - ② Mark the current vertex as visited.
 - ③ Select the unvisited vertex with smallest tentative distance and set it as the current vertex.

Dijkstra's Algorithm – Correctness

Let S be the set of marked vertices, u the starting vertex, and $d(v)$ the calculated distance from u to v , then the correctness of Dijkstra's algorithm can be stated as:

Lemma

For every vertex $v \in S$ $d(v)$ is the shortest path from u to v .

Dijkstra's Algorithm – Correctness

Proof. By induction on $|S|$.

- **Base Case.** Easy, $|S| = 1$ implies $S = \{u\}$, and the distance is 0. Alternatively $|S| = 0$ is trivial.
- **Induction Hypothesis.** Assume the theorem is true for $|S| = k \geq 1$.
- **Inductive Step.**
 - Let v be the next vertex added to S , and w be the neighbouring path vertex according to the algorithm.
 - By the assumption, $d(w)$ is minimum.
 - Let P , for contradiction, be the shortest path from u to v .
 - Let xy be the first edge in P that leaves S .
 - Then the length of P is at least $d(x) + w(x, y)$, which is the tentative distance to y .
 - This distance can't be smaller than the tentative distance to v , otherwise we would've chosen it instead of v to add next.
 - Therefore the length of P is at least $d(v)$. □

- We can keep an array of predecessors that gets updated when we update the distance – this will give us the shortest path tree.
- Where was the dynamic programming?
 - When we select whether to update the distance or keep the existing one.
 - If $dist[v]$ is the distance from u to v , then the core of the algorithm is repeatedly performing
$$dist[v] = \min\{dist[v], dist[current] + w(current, v)\}.$$
- Note there's also a greedy selection of the next vertex to process.
- Where's the problem with negative edges?

Dealing with Negative Edges - Bellman-Ford Algorithm

- Dijkstra's algorithm ran into trouble with negative edges because we never look at marked vertices again.
- The Bellman-Ford algorithm can handle negative edges – it just looks at everything again!
- Doesn't work with negative *cycles* though – but it can detect them, and they pose a semantic problem anyway.
- Historical note: Both algorithms were invented at roughly the same time, Dijkstra's in 1956 (but not published until 1959), Bellman-(Moore)-Ford in '58, '57 and '56 respectively.

- ① Set things up as with Dijkstra's.
- ② For 1 to the number of vertices
 - ① For each edge uv
 - ① If using uv improves the distances, use it and update as needed.
- ③ Check for negative cycles:
 - ① For each edge uv
 - ① If uv improves the distances, there's a negative cycle.

Bellman-Ford Algorithm - Correctness

Formally, we need a slightly more precise statement:

Lemma

After i iterations of the outer for loop:

- *If $d(v) \neq \infty$, there is a (known) path from the start vertex u to vertex v with distance $d(v)$.*
- *If there is a path from u to v with at most i edges, $d(v)$ is the length of the shortest path with at most i edges.*

Bellman-Ford Algorithm - Correctness

Proof. By induction on i :

- **Base Case:** $i = 0$, nothing has happened yet, trivial!
- **Inductive Assumption:** Assume that at step i we have the shortest paths with at most i edges.
- **Inductive Step:**
 - (First part) When $d(v)$ is updated to $d(r) + w(r, v)$, there is a path from u to r to v with weight $d(v)$.
 - (Second part) Before updating, we have by assumption the distances of all the shortest paths with at most i edges computed (if they exist).
 - Then at each possible update, we compare the shortest path to a neighbour plus an edge, with the shortest path on at most i vertices, and keep the minimum.
 - So we have a path on $i + 1$ edges that is shorter than the i edge path (and we pick the smallest of all these), or we keep the i edge path – thus the path must be the shortest path on $i + 1$ edges – there are no other possible ways to get it. \square

Dijkstra's vs. Bellman-Ford – Complexity

- Dijkstra's:
 - $O(n^2)$ time.
 - $O(m + n \log n)$ time if we use a priority queue to help pick the next vertex.
 - $O(n)$ space.
- Bellman-Ford:
 - $O(nm)$ time – can't really get around this.
 - $O(n)$ space.

Connectivity

- When using graphs to model things, we usually want to know if they're all in one piece.
- We've used the word "connected" before.
- A graph is *connected* if there's a way to get from any vertex to any other vertex via the edges.
- A pretty useful property if you're running a computer network, or transport network...

- Testing this property is actually easy.
- Just do a traversal!
- If we visit all the vertices, the graph must be connected. If we don't, it's not.

What if it's not connected?

- If it turns out that the graph is not connected, we might be interested in finding out what bits (subgraphs) are.
- These are called *connected components*.
- Finding them is almost as easy as testing connectivity:
 - ① While not all vertices have been visited:
 - ① Pick an unvisited vertex.
 - ② Start a traversal there.
 - ③ All the vertices you can reach from there form a component, record this and mark them as visited.

Strongly Connected Components

In directed graphs, things get more complicated:

- In an undirected graph, if a is adjacent to b then b is also adjacent to a .
- Not necessarily true in directed graphs.
- In undirected graphs, every vertex in a connected component is reachable from every other vertex in the same component.
- Not necessarily true in directed graphs.

Strongly Connected Components

The directed graph analogy to this reachability based connectivity is *strong connectivity*.

- Two vertices are *strongly connected* if there is a path from one to the other, and a path back again.
- This gives a *strongly connected component*: a subgraph where every pair of vertices is strongly connected.

How do we find these?

Strongly Connected Components

We can just run a bunch of reachability queries.

- Parallelises well.
- $O(n^2(n + m))$ without clever optimisations (keeping track of all the vertices visited in the query).
- Can get it to $O((n + m) \log n)$ expected sequential time if we make the approach complicated enough (uses divide and conquer).

Strongly Connected Components

We can do it faster overall (but without the boost from parallelism) using *Tarjan's Algorithm*.

- Depth first traversal based.
- Each vertex is assigned an *index* and a *low* value.
 - *index* keeps track of where it is in the traversal.
 - *low* tracks the lowest index reachable from this vertex.
- If *low* is less than *index*, there's a path back to something earlier in the traversal (and of course one forward - the traversal itself), so all those vertices must be in a strongly connected component together.

Tarjan's SSC Algorithm 1

```
set<set<vertex>> tarjan(directed_graph<vertex> d) {  
    int index = 0;  
    stack s;  
    for (each vertex v in d) {  
        if (index[v] is undefined) {  
            components.add(strongconnect(v));  
        }  
    }  
}
```

Tarjan's SSC Algorithm 2

```
set<vertex> strongconnect(vertex v) {
    index[v] = index; low[v] = index; index = index + 1;
    s.push(v); on_stack[v] = true;

    for (each edge (v, w) in d) {
        if (index[w] is undefined) {
            strongconnect(w);
            low[v] = min(low[v], low[w])
        } else if (w.onStack){
            low[v] = min(low[v], index[w])
        }
    }

    if (low[v] == index[v]) {
        set<vertex> new_component;
        do {
            w = s.pop(); on_stack[w] = false;
            new_component.add(w);
        } while (w != v)
        return new_component;
    }
}
```

- Time: $O(n + m)$.
- Space: $O(n)$ (what extra stuff was kept track of in the algorithm?).

Back in undirected land, a similar algorithm can be used to find the equivalent of strongly connected components.

- Which vertices (or edges) can we remove that will increase the number of connected components?
- These vertices are called *articulation* points (the edges are called *bridges*).
- A connected component that has no articulation points is called *biconnected*.

Finding Articulation Points

- We can find articulation points using an augmented depth first traversal.
 - We need to keep track of the depth of a vertex in the traversal, and
 - the lowest depths of the neighbours of each vertex and all its descendants in the DFS tree (the “lowpoint”).
- If a vertex's depth is less than or equal to the lowpoint of any of its neighbours, it's an articulation point!
- The lowpoint is actually just a way of tracking whether there's at least two ways to get to a vertex – vertices in biconnected components all have the same lowpoint.
- The root vertex is a special case – if it has at least two children, it's an articulation point.

```
FindArticulationPoints(i, d)
    visited[i] = true, depth[i] = d
    low[i] = d, child_count = 0
    is_articulation = false
    for (each neighbour n of i)
        if (not visited[n])
            parent[n] = i
            FindArticulationPoints(n, d+1)
            child_count = child_count + 1
            if (low[n] >= depth[i])
                is_articulation = true
            low[i] = min{low[i], low[n]}
        else if (n != parent[i])
            low[i] = min{low[i], depth[n]}
    if ((i has a parent and is_articulation) or
        (i has no parent and child_count > 1))
        i is an articulation point
```

- Time: $O(n + m)$ – it's just a depth-first traversal.
- Space: $O(n)$
- Where n and m are the number of vertices and edges respectively.
- Can be modified to retrieve the biconnected components with the same complexity.
- Can be done in parallel in $O(\log n)$ time! (But with $O(n + m)$ processors...)

Dependencies in Directed Graphs

If a directed graph has no directed cycles, we called it a *Directed Acyclic Graph*, or DAG.

How many strongly connected components does this have?

DAGs are like the trees of the directed graph world, and as such have many applications:

- Anything where dependencies are important (anyone running a *NIX distro will understand this one).
- Task scheduling.
- Compilation.
- Computing information in a spreadsheet.
- Modelling multi-stage processes.

For all of these, it's pretty clear we need to know what order to do things in.

This ordering is called a *topological ordering*. It may not be unique. (What if it is?)

To find a topological ordering, we perform a *topological sort*.

```
list<vertex> topo_order;
set<vertex> s = {all vertices with no incoming edges};
while (!s.empty()) {
    v = s.remove_something_from_a_set();
    topo_order.append(v);
    for (each neighbour u of v) {
        remove (v,u) from graph;
        if (u has no incoming edges) add u to s;
    }
    if (any edges left in graph) return error;
else return topo_order;
```

Depth First Topological Sort

```
list<vertex> topo_order;

while (not all vertices permanently marked) {
    v = next unmarked vertex;
    visit(v)
}

visit(vertex v) {
    if (v is permanently marked) return;
    if (v is temporarily marked) return error;
    mark v temporarily;
    for (each (v,u) in edges) {
        visit(u);
    }
    clear temporary mark from v;
    permanently mark v;
    topo_order.append_front(v);
}
```

- Both linear time algorithms – $O(n + m)$.
- Both use $O(n)$ extra space (can do some tricks here to get $O(1)$ for Kahn's algorithm at the cost of increasing time to $O((n + m) \log n)$).
- Can find a topological (a different way) in $O(\log^2 n)$ time ... if we have a polynomial number of processors.