

31251 – Data Structures and Algorithms

Week 11

Luke Mathieson

Everyone's gonna love this:

- Moving from Algorithmic Complexity to Computational Complexity.
- **P**.
- **NP**.
- **NP**-hardness and **NP**-completeness.

So far we've seen a variety of algorithms for problems.

- We can compare them via their asymptotic running time.
- What about the problems themselves?
- An algorithm is a solution for a problem, but it may not be the best one.
- How can we get some idea of whether we're doing well with our algorithms?

We can talk about the complexity of *problems* too, but we need some formalisations...

Formalising Computational Problems

- A (basic) computational problem consists of two parts:
 - The input.
 - The output.
- Both should be specified *precisely*.

For example, we can express sorting an array as:

Sort

Input: An array A of length n .

Output: An array A' with the same elements as A where $\forall 1 \leq i \leq j \leq n$ we have $A'[i] \leq A'[j]$.

Formalising Computational Problems

The same computational problem can come in several varieties:

- Optimisation - we want to find the best solution according to some optimisation criterion.
- Search - we want to find any solution that solves the problem.
- Decision - we just want to know *if* there is a solution.

For a variety of reasons, we're going to be most interested in *decision* problems. (Are these always equivalent to the others?)

Formalising Computational Problems

So the decision version of sorting would look like:

`SORT`

Input: An array A of length n , two integers k, p .

Question: Is element $A[k]$ at position p when A is put into ascending order?

For those who like formal languages (ha!) a decision problem can be equivalently thought of as the language containing the strings encoding the YES instances.

The Complexity of a Problem

Now that we have at least some notion of a formal problem, we define what we mean by its complexity:

Definition

The complexity of a problem is the complexity of the (asymptotically) most efficient algorithm that correctly solves the problem.

- Most of the time, when we say “most efficient”, we mean fastest.
- Any algorithm that solves the problem at least gives an upper bound – but we also need a lower bound to be precise.

“Most Efficient” on what?

To make this definition well-formed, we need to specify how we're going to measure efficiency.

- What is the machine model that we're going to use?
- RAM vs. Turing Machine? (How many tapes?)
- What's a “unitary” operation?

We're going to take Turing Machines as our basic model. (Why?)

So using this model, we can say things like:

- SORT has complexity $\Theta(n \log n)$.
- DEPTH FIRST SEARCH has complexity $\mathcal{O}(|V| + |E|)$.
- MINIMUM SPANNING TREE has complexity $\mathcal{O}(|E| \cdot \log |V|)$.
- SHORTEST PATH(S) has complexity $\mathcal{O}(|E| + |V| \log |V|)$.

These are fast, practical algorithms - we would like to group these together...

Definition (Polynomial-time or “P”)

Given a decision problem Π , we say that Π is polynomial-time solvable, or $\Pi \in \mathbf{P}$, if and only if there exists a Turing Machine (\approx algorithm) \mathcal{A} such that for each instance I of Π the Turing Machine \mathcal{A} halts and correctly answers YES or NO in time bounded by $\mathcal{O}(|I|^c)$ for a fixed $c \in \mathbb{N}$.

Alternatively, if we think of Π as a language, \mathcal{A} decides $I \in \Pi$ in polynomial time.

The Class **P** & the Cobham Edmonds Thesis

Cobham Edmonds Thesis

Computational problems can be feasibly computed on some computational device only if they can be computed in polynomial time.

- This is the basic rule-of-thumb that standard complexity theory works with.
- Why is this a reasonable thing to say? (*Is it reasonable?*)

Excursion 1 - Measuring Things

- Notice that in the definition of \mathbf{P} , we talk about the size of the instance $|I|$, what does this mean?
- The formal definition is that $|I|$ is the length of a reasonable encoding of I over some alphabet of symbols Σ (the input alphabet of the Turing Machine).
- We can reasonably assume that $\Sigma = \{0, 1\}$, so the encoding is in binary. (Why?)
- But what does it mean when we start using $|V|$ or $|E|$, or other measures for the complexity?
 - These measures are called *proxy* measures.
 - Note given a graph, we can encode it in at most $\mathcal{O}(|V|^2 \log |V|)$ bits. (Why?)
 - So $|V|$ is within a polynomial factor of $|I|$, so it's "okay" to use it as a proxy.

Is Everything in **P**?

- So 'everything' we've seen so far is in **P**.
- Is *everything* in **P**?
- Of course not! We can easily make up problems that aren't, some things aren't even decidable!
- What about things that we might want to solve reasonably, but can't quite seem to get a polynomial-time algorithm? Can we say anything here?

Excursion 2 - Non-determinism

- You might remember a brief mention of non-determinism several weeks ago.
- Turing Machines can be non-deterministic too.
- What does this mean?
 - Non-determinism is a piece of theoretical “magic”.
 - A non-deterministic machine can guess cleverly to solve a problem (but we still have to check the solution).
 - Slightly more formally, a non-deterministic machine, on a given input, when presented with more than one available transition, will follow the path that leads to an accepting configuration.
 - In other words, if there's a way it can say YES, it will - this doesn't guarantee it's correct though!.

- The class **P** is actually the class of problems that can be solved in polynomial time by a **deterministic** Turing Machine.
- What happens if we switch deterministic for non-deterministic?
- We get **NP**!

Non-deterministic Polynomial Time, or **NP**

Definition (**NP**)

Given a decision problem Π , we say that Π is non-deterministically polynomial-time solvable, or $\Pi \in \mathbf{NP}$, if and only if there exists a *non-deterministic* Turing Machine (\approx algorithm) \mathcal{A} such that for each instance I of Π the Turing Machine \mathcal{A} halts and correctly answers YES or NO in time bounded by $\mathcal{O}(|I|^c)$ for a fixed $c \in \mathbb{N}$.

- Obviously $\mathbf{P} \subseteq \mathbf{NP}$ - just don't do any guessing.
- Do we get any more power from non-determinism (i.e. is $\mathbf{P} \subset \mathbf{NP}$?)
- We'll come to that later.

An Alternate Characterisation of **NP**

NP can also be thought of as the class of problems that we can verify in polynomial-time:

Definition (Polynomial-time verifiable.)

Given a decision problem Π , we say that Π is polynomial-time verifiable, if and only if there exists a deterministic Turing Machine (\equiv algorithm) \mathcal{A} such that for each instance I of Π , with a witness W that I is a YES instance where $|W| \leq |I|^k$ for some $k \in \mathbb{N}$, the Turing Machine \mathcal{A} halts and correctly answers whether the witness is correct in time bounded by $\mathcal{O}(|I|^c)$ for a fixed $c \in \mathbb{N}$.

That is, given a solution we can check if it's right quickly. It turns out that this is the same as **NP**

Why are these two definitions equivalent?

A *third* definition: **NP** is the class characterised by existential second-order logic (Fagin's Theorem).

So, for example,

- SORT is in **NP** - check the final array in linear time.
- MST is in **NP** - add up the weights of the solution and check (note that we're always talking about the decision version).

But there are things that are in **NP** for which we don't know any polynomial-time algorithm.

It seems like **NP** contains a 'lot more' than **P** - many problems that have no obvious **P**-time algorithm are in **NP**.

- No one knows whether $\mathbf{P} = \mathbf{NP}$ or not though.
- The general appraisal is that $\mathbf{P} \neq \mathbf{NP}$, but there is notable disagreement.

If we can't implement a non-deterministic algorithm, what is the point of **NP**?

- **NP** gives us the basis of a way to show that a problem has no polynomial-time algorithm unless $\mathbf{P} = \mathbf{NP}$.

Excursion 3 - Reducibility

Before we can say how hard things are (when we don't have an algorithm), we need to be able to compare them. We do this via *polynomial-time many-one mapping reductions*, AKA ptime reductions, or Karp reductions:

Definition

Given two decision problems A and B , we say that A is (ptime many-one) reducible to B (sometimes written $A \leq_m B$) if and only if there exists a deterministic algorithm that takes an instance I_A of A and maps it to an instance I_B of B such that I_A is a YES-instance if and only if I_B is a YES-instance, where the algorithm runs in time $\mathcal{O}(|I_A|)$.

Excursion 3 - Reducibility

- So if we could solve B in \mathbf{P} -time, we could solve A in \mathbf{P} -time (start with I_A , map to I_B , solve I_B).
- If B is in \mathbf{P} , and $A \leq_m B$, then A must also be in \mathbf{P} .
- Note that as the running time is polynomial, $|I_B| \leq |I_A|^c$.
- So then B is “at least as hard as” A , up to some polynomial factor.

Excursion 3 - Reducibility

- So reductions give us an ordering on decision problems.
- So what happens when start looking for reductions between problems in **NP**?
- It turns out that there's a group of *really hard* problems that cluster together at the 'end' of **NP**.
- These are the **NP**-complete problems.

NP-hardness and NP-completeness

- 1 A problem Π is **NP**-hard if for every $\Pi' \in \mathbf{NP}$ there exists a Karp reduction such that $\Pi' \leq_m \Pi$.
- 2 If Π is also in **NP** (note we did not make this assumption above), then Π is **NP**-complete.

If a problem A is **NP**-hard and $A \leq_m B$, then B is also **NP** – *hard*.
(Why?)

So the **NP**-complete problems are the ‘hardest’ problems in **NP** - if there’s a polynomial time algorithm for any of them, then $\mathbf{P} = \mathbf{NP}$.

The Cook-Levin Theorem, or, SAT is **NP**-complete

The first **NP**-complete problem is SATISFIABILITY (SAT):

SAT

Instance: A set of variables V and a set of clauses C over the literals of V (i.e. a CNF formula).

Question: Is there an assignment of TRUE and FALSE to V such that the entire formula evaluates to TRUE?

Theorem (Cook-Levin Theorem)

SAT is **NP**-complete.

- Then from SAT, we can reduce to thousands of other problems.
- Some you may see or hear of:
 - VERTEX COVER.
 - HITTING SET.
 - DOMINATING SET.
 - INDEPENDENT SET.
 - CLIQUE.
 - TRAVELLING SALESMAN PROBLEM.
 - GRAPH COLOURING.

What do we do with all this?

- So now we know that some things are **NP**-complete. What next?
- This gives strong evidence that looking for a polynomial-time algorithm is a waste of time.
 - An efficient algorithm for *any* **NP**-complete problem would give an algorithm for all of them, and no-one has found one yet.
- Even worse, it turns out that lots of problems are **NP**-complete.
 - Sometimes it seems like there's only about 6 problems in **P**.
- So we have to develop methods for coping with intractability.

Methods for Coping with Intractability

- Approximate - maybe we're happy with not-quite-the-best solutions.
- Randomize - maybe we can get lucky most of the time.
- Develop heuristics and metaheuristics - use what seems to work in practice.
- Relax what it means to be tractable - sometimes exponential is better than polynomial.
- Build a different type of computer - will at least get you research grants.